

Eashwar Subramanian

Data Analyst | SQL, Python, Tableau / Power BI | Data Quality, Reporting & Analytics

Melbourne, VIC | eashwars2001@gmail.com | +61 450 332 782

LinkedIn: linkedin.com/in/eashwar-subramanian-02461a1b3 | **GitHub:** github.com/Eashwar-Subramanian | **Portfolio:** eashwar-subramanian.github.io/Portfolio

Work rights: Full-time work rights in Australia; Temporary Graduate visa (subclass 485) valid until February 2029

SUMMARY

Data Analyst with hands-on experience across data pipelines, dashboards, reporting workflows, and quality-focused analytics. Experienced in collecting, cleaning, validating, and structuring structured and unstructured data using SQL, Python, Tableau, Power BI, PostgreSQL, and MySQL. Strong in reporting logic, metric validation, deduplication, QA checks, and documentation, with a focus on building reproducible workflows that support stakeholder-ready dashboards, briefs, and decision-making. Comfortable translating unclear business questions into structured analytical workflows and refining outputs through iteration, validation, and feedback.

CORE SKILLS

SQL & Databases: SQL, PostgreSQL, MySQL, SQL Server, joins, CTEs, views, temp tables
BI & Reporting: Tableau, Power BI, Excel, KPI reporting, dashboard design, drilldowns, source-to-dashboard validation
Python & Analytics: Python, pandas, NumPy, BeautifulSoup, forecasting (SARIMA / SARIMAX), segmentation (RFM, K-Means), evaluation (RAGAS)
Data Quality & Governance: data cleaning, deduplication, metric validation, audit-style QA checks, metadata tagging, audit logging, Row Level Security (RLS)
Tools & Delivery: Git, Docker, Pytest, Jira, Confluence, Bitbucket
Cloud / Application Exposure: AWS (S3, Lambda, EC2), FastAPI, Flask, Folium
Other: reproducible workflows, taxonomy design, schema mapping, monitoring / evaluation, documentation

PROFESSIONAL EXPERIENCE

Cultural Infusion | Melbourne, Australia
Research Assistant (Data Analytics) | Mar 2026 – Present

- Develop and refine the initial analytical framework for a proposed study on cultural and demographic predictors of disinformation vulnerability in Cyprus and the United Kingdom, with emphasis on policy relevance and defensible outputs.
- Define a three-layer methodology: identity/demographic cluster construction, vulnerability indicator scorecards, and narrative targeting overlays linking active narratives to potentially vulnerable profiles.
- Scope and assess comparative data inputs across public survey/population sources (e.g., ESS, Eurobarometer, Ofcom) and public narrative sources (including GDELT and other public channels), prioritising comparability and grounding.
- Set validation checkpoints across data collection, clustering, narrative mapping, and policy interpretation to reduce over-claiming and keep outputs aligned with the methodology.
- Draft output concepts (cluster cards, matrix/heatmap visuals, short methodology note) and a concise presentation structure for stakeholder review.

Data Science Intern (Project Atlas – Public Intelligence Monitoring Pipeline) | Jan 2026 – Mar 2026

- Delivered an evidence-based public intelligence pipeline that converts news and public narrative sources into structured, decision-ready briefs with direct source traceability.
- Orchestrated multi-source ingestion and evidence extraction workflows spanning GDELT news retrieval, public/social narrative retrieval, article text extraction, and normalized topic-level storage.
- Applied staged relevance judgement (retain/prune) and story-level clustering to reduce noise, consolidate overlapping coverage, and support more stable downstream analysis.
- Enabled structured signal extraction and delta-aware brief generation so recurring runs could surface net-new developments rather than repeating previously seen coverage.
- Configured scheduled watchlist monitoring with Teams alerts sent only when new items were detected, helping reduce alert fatigue in recurring tracking.
- Contributed to a web-backed, multi-user workflow supporting manual runs, scheduled watchlists, archive viewing, topic ownership, and optional PostgreSQL-backed application features.

Solara Health | Melbourne, Australia
Data Science Intern | Jul 2025 – Nov 2025

- Built ingestion pipelines for a healthcare content library covering PDF and HTML sources, applying metadata governance and SHA-256 deduplication to improve dataset reliability and traceability.
- Automated download, parsing, and transformation workflows with retries, content-type handling, and standardized text outputs to stabilize downstream retrieval datasets.
- Designed PostgreSQL data structures for healthcare content, metadata, and embeddings, and exposed curated datasets through a FastAPI service with basic automated tests for reliability.
- Implemented tenant-aware access controls using PostgreSQL Row Level Security (RLS) to enforce segregation across multi-hospital deployments.
- Developed QA and audit tooling including flag logging, CSV exports, citation audit checks, labelled evaluation datasets, and repeatable evaluation reporting using RAGAS and custom checks.
- Collaborated in an Agile team environment using Jira, Confluence, and Bitbucket, contributing documentation, validation rules, delivery planning, and iterative improvements across sprint cycles.

Preplnsta Pvt Ltd (Remote) | India
Data Analyst Intern | Dec 2023 – Feb 2024

- Built and delivered stakeholder dashboards in Tableau and Excel, validating source tables, transformations, and dashboard calculations to improve reporting consistency.
- Optimized multi-table SQL queries for recurring executive reporting and ad hoc reporting requests.
- Automated recurring reporting inputs in Python using BeautifulSoup-based extraction, improving repeatability and reducing manual effort in reporting workflows.

PROJECTS

Solara Healthcare RAG System | Jul 2025 – Nov 2025

- Contributed to a team-built healthcare RAG platform for Australian oncology support, covering document ingestion, retrieval, evaluation, safety, and multi-tenant access control.
- Worked on evaluation and monitoring workflows using RAGAS and related quality checks; the final project report records faithfulness of 69.4%, answer relevancy of 54.3%, context precision of 92.7%, and context recall of 91.2%.
- Contributed to a production-style architecture using semantic chunking, cross-encoder reranking, PostgreSQL with pgvector, Docker, FastAPI, and tenant-scoped retrieval controls.

Tools: Python, FastAPI, PostgreSQL, pgvector, Docker, RAGAS, LlamaIndex, Sentence-Transformers, AWS

Sales & Customer Performance Dashboard | Tableau

GitHub: github.com/Eashwar-Subramanian/sales-customer-performance-bi-tableau

- Built two interactive Tableau dashboards focused on sales performance and customer performance, covering current-year vs previous-year KPIs, sub-category trends, weekly results, and customer profitability views.
- Implemented parameter-driven CY / PY calculations, drilldowns, cross-filtering, dashboard navigation, and a show / hide filter panel to support cleaner self-service exploration.
- Structured the build around a reproducible BI workflow with mockups, technical notes, data dictionary documentation, and packaged Tableau assets.

Tools: Tableau, CSV data modelling, calculated fields, parameters, dashboard actions

Sales Analytics Dashboard | Power BI, MySQL

GitHub: github.com/Eashwar-Subramanian/sales-analytics-dashboard-powerbi

- Built a Power BI dashboard connected to a MySQL dataset to track performance across time, region, and product category.
- Created reporting-ready views and consistent metric definitions to support recurring performance reporting and stakeholder review.
- Used KPI summaries, trend views, and drilldowns to help users interpret performance changes quickly and identify follow-up areas for analysis.

Tools: Power BI, MySQL, KPI reporting, dashboard modelling

Australian Retail Customer Segmentation | Python, SQL, Power BI

GitHub: github.com/Eashwar-Subramanian/retail-customer-segmentation-australia

- Built an end-to-end segmentation workflow covering transaction-integrity checks, data cleaning, customer-level feature engineering, RFM analysis, K-Means clustering, and dashboard-ready outputs.
- Segmented 788 customers and identified a 106-customer high-value cohort contributing about 15.9% of revenue, then translated cluster behavior into retention, cross-sell, and win-back recommendations.
- Validated order integrity before modelling, finding widespread inconsistency in reused order numbers and repairing the workflow so segmentation and reporting were based on more reliable customer-transaction units.

Tools: Python, pandas, NumPy, K-Means, RFM analysis, Power BI, Jupyter Notebook

Climate Policy Forecasting Dashboard | Python, Flask, SARIMAX, Folium

GitHub: github.com/Eashwar-Subramanian/climate-policy-forecasting-dashboard

- Built a Flask dashboard that generates on-demand weekly SARIMAX forecasts for Australian locations and returns forecast outputs through JSON responses.
- Added weekly resampling, interpolation, and MAE / RMSE validation logic for model checking where sufficient historical data was available.
- Integrated a temporal Folium map and location-based exploration layer to make geographically distributed forecast outputs easier to interpret.

Tools: Python, Flask, stats models, pandas, scikit-learn metrics, Folium, matplotlib

Real Estate Data Cleaning in SQL | MySQL

GitHub: github.com/Eashwar-Subramanian/real-estate-data-cleaning-mysql

- Built a reusable SQL cleaning workflow covering date standardization, missing-address filling, address splitting, categorical normalization, duplicate identification, and column cleanup.
- Produced cleaner, analysis-ready tables while documenting the transformation logic step by step for repeatability and review.

Tools: MySQL, SQL data cleaning, transformation logic, duplicate handling

COVID-19 Analysis | SQL, Tableau Prep

GitHub: github.com/Eashwar-Subramanian/covid-data-analysis-sql-tableau

- Wrote SQL queries to analyze COVID-19 cases, deaths, population-level infection rates, and vaccination trends using joins,

aggregations, window functions, CTEs, and views.

- Prepared Tableau-ready outputs from SQL workflows to support downstream visual analysis of global and country-level patterns.

Tools: SQL, MySQL, joins, window functions, CTEs, Tableau Prep outputs

Climate-Driven Urban Growth Analysis | R

GitHub: github.com/Eashwar-Subramanian/climate-urban-growth-analysis-r

- Integrated meteorological and city-level population datasets to examine whether rainfall, temperature, and humidity signals related to population growth across Australian cities from 2011 to 2017.
- Performed missing value handling, outlier treatment, location matching, correlation analysis, and visual exploration using reproducible R workflows.

Tools: R, dplyr, tidyr, ggplot2, correlation analysis

E-Commerce Shipping Analysis | R

GitHub: github.com/Eashwar-Subramanian/ecommerce-shipping-analysis-r

- Conducted exploratory analysis on a 10,999-row e-commerce shipping dataset to understand which operational factors were associated with the delivery-timeliness flag.
- Analyzed relationships across shipment mode, warehouse block, product importance, discount level, package weight, and customer interaction signals to surface operational patterns for review.

Tools: R, exploratory data analysis, visualization, dataset profiling

Movies Correlation Analysis | Python

GitHub: github.com/Eashwar-Subramanian/movies-correlation-analysis

- Performed exploratory data analysis on movie attributes with a focus on relationships between budget, gross revenue, and other numeric variables.
- Cleaned and prepared analysis-ready columns, then used correlation analysis and visual exploration to highlight major relationships in the dataset.

Tools: Python, pandas, NumPy, Jupyter Notebook, exploratory analysis

Cloud Music Subscription App | Java, AWS

GitHub: github.com/Eashwar-Subramanian/cloud-music-subscription-app

- Built a coursework full-stack music subscription application with Java servlet components supporting login, registration, search, and subscribe / unsubscribe flows.
- Integrated AWS components including DynamoDB-backed persistence and S3-based image handling utilities as part of the application workflow.

Tools: Java, Servlets, AWS S3, DynamoDB, Lambda, HTML

EDUCATION

Master of Data Science

Royal Melbourne Institute of Technology | Melbourne, Australia
Dec 2025 | GPA: 3.6/4.0

Bachelor of Electronics & Communication Engineering

Rajagiri School of Engineering and Technology | Kerala, India
Aug 2023 | First Class Distinction | GPA: 8.8/10